



Masterarbeit

LLM-based Academic Writing with Fine-Grained Citations

Eberhard Karls Universität Tübingen
Mathematisch-Naturwissenschaftliche Fakultät
Wilhelm-Schickard-Institut für Informatik
Lernbasierte Computer Vision
Niklas Munkes, niklas.munkes@student.uni-tuebingen.de, 2025

Bearbeitungszeitraum: 14.04.-14.10.2025

Gutachter:	Prof. Dr. Andreas Geiger, Universität Tübingen
Zweitgutachter:	Prof. Dr. Carsten Eickhoff, Universität Tübingen
Betreuer:	Haoyu He, Universität Tübingen

Selbstständigkeitserklärung

Hiermit versichere ich, dass ich die vorliegende Masterarbeit selbständig und nur mit den angegebenen Hilfsmitteln angefertigt habe und dass alle Stellen, die dem Wortlaut oder dem Sinne nach anderen Werken entnommen sind, durch Angaben von Quellen als Entlehnung kenntlich gemacht worden sind. Diese Masterarbeit wurde in gleicher oder ähnlicher Form in keinem anderen Studiengang als Prüfungsleistung vorgelegt.

Niklas Munkes (Matrikelnummer 4269436), October 2, 2025

Abstract

TODO

Contents

1. Introduction	9
2. Related Work	11
2.1. LLM-based Generation	11
2.2. RAG for Scientific Writing	11
3. Methods	13
3.1. Retrieval-Augmented Generation	13
3.2. Dataset	13
3.3. ScholarCopilot	15
3.4. Passage Retrieval	16
4. Experiments	21
4.1. Evaluation Methodology	21
4.2. Single-Step Retrieval	23
4.3. Generation Quality	28
4.4. Additional Experiments	29
5. Limitations and Future Work	31
6. Conclusion	33
A. Appendix	35
A.1. Hyperparameters	35
A.2. Passage Retrieval Instruction	35
A.3. Generation Evaluation Instruction	35
A.4. Complete Example of Single-Step Retrieval	35
A.5. "False Positive" Example of LLM-based Generation Quality Evaluation	37

1. Introduction

TODO

citations build trust [DOW⁺24]

Retrieval-Augmented Generation (RAG, [LPP⁺20]) models are designed to generate text that is grounded in real-world factual knowledge. They have been widely adopted for knowledge-intensive NLP tasks [SQK⁺25] such as long-form question answering [XWL⁺24], scientific question answering [LOS⁺23] and question synthesis [XLS⁺24], generation with sentence-level citations [ZBL⁺24, XWL⁺24] and academic text generation [AHS⁺24, WMN⁺25]. RAG models usually consist of a generator, a pre-trained sequence-to-sequence (seq2seq) model (e.g. a Transformer-based [VSP⁺17] decoder), and a retriever with access to a retrieval corpus, which may be implemented as a dense vector index (e.g. a HNSW [MY18] index) [LPP⁺20]. The RAG architecture is briefly introduced in section 3.1. Such models combine pre-trained parametric memory with non-parametric (retrieval-based) memory [LPP⁺20], which has a number of benefits: while neural language models are able to store a substantial amount of world knowledge in their parameters [RRS20], this knowledge can't easily be modified, expanded or directly be used to gain insight into their decision-making process [LLG⁺19]. By using an external retrieval corpus, the knowledge RAG models use to inform their generation can directly be inspected and interpreted [LPP⁺20]. Their TODO

TODO

TODO

2. Related Work

TODO

2.1. LLM-based Generation

[STB25] TODO

2.2. RAG for Scientific Writing

Previous work on text generation with citations has recently transitioned from treating retrieval and generation as separate processing stages (e.g. [SHC⁺24]) to alternating between the two (e.g. [XWL⁺24], [AHS⁺24]). *Attribute First, then Generate* [SHC⁺24] initially retrieves a set of relevant sub-sentence spans which are then grouped into coherent clusters which in turn are used to inform an iterative sentence-by-sentence generation step [SHC⁺24]. *ReClaim* [XWL⁺24], more specifically the *ReClaim with Interleaving Generation* method, is used to generate an output sequence $O = \{r_1, c_1, r_2, c_2, \dots, r_n, c_n\}$ of alternating references and claims where claim c_i is the portion of the model's response that was generated based on reference r_i . Claim and reference generation is performed by separate LLMs specifically trained for their corresponding task [XWL⁺24]. In contrast, *OpenScholar* [AHS⁺24] first retrieves and reranks a set of papers P relevant to the given query. These papers inform a subsequent generation step that produces a generated response r_i [AHS⁺24]. This response is refined through iterative *self-feedback* where during each iteration of the feedback loop $F = f_1, f_2, \dots, f_T$, the generator model generates sentences f_j that describe potential improvements to the response and may also prompt the retrieval pipeline to retrieve additional papers, should r_i require more information to generate an improved response r_{i+1} [AHS⁺24]. Depending on the size of P and the choice of T , OpenScholar still performs retrieval and generation largely separately.

ScholarCopilot [WMN⁺25] is a RAG system developed for generating academic-style text with accurate citations. It is intended as a tool that assists the user with academic paper writing by providing iterative text generation with automated citation retrieval [WMN⁺25]. ScholarCopilot also enables user interaction during generation [WMN⁺25] by allowing the user to rewrite the generated text or forcefully

Chapter 2. Related Work

trigger citation retrieval between the consecutive generation steps but this feature is beyond the scope of this thesis and is thus not utilized for the experiments. ...

3. Methods

TODO

3.1. Retrieval-Augmented Generation

The RAG architecture [LPP⁺20] is comprised of a retriever p_η consisting of a query encoder q and a retrieval corpus d , and a generator p_θ (see Figure 3.1). Given a query x , the retriever $p_\eta(z|x)$ (with parameters η) computes a distribution over the text documents z in the retrieval corpus d according to x [LPP⁺20]. This distribution is used to obtain the top- k most relevant documents with respect to the query. The generator $p_\theta(y_i|x, z, y_{1:i-1})$, parametrized by θ , then predicts each next token y_i based on the previous tokens $y_{1:i-1}$, the given query x and a document z from the top- k retrieved documents (according to p_η) [LPP⁺20]. Lewis et al. base their retriever on the Dense Passage Retriever [KOM⁺20] and use a pre-trained BART-large [LLG⁺19] as the generator.

3.2. Dataset

This thesis evaluates the retrieval performance and generation quality of ScholarCopilot (SC, see section 3.3) with and without the passage retrieval module (PR, see section 3.4). ScholarCopilot retains its original retrieval corpus, while the retrieval corpus for the PR module uses the HDT dataset [HFB⁺24], since it requires fulltext data and the SC retrieval corpus only consists of abstracts. The PR module constructs the passages for each individual document during inference. The datasets used for evaluation are constructed from the HDT dataset as well.

The datasets for retrieval evaluation (section 4.2) each use a subset of the HDT dataset. Each test sample in the dataset consists of a citing text span and the corresponding cited document. All citing spans are extracted from the documents in the HDT dataset, provided the corresponding cited document appears in the ScholarCopilot retrieval corpus. This results in a total of 5.16M citing contexts. Following the evaluation setup in [WMN⁺25], all tokens after the citation are removed. This emulates the context the model would receive during generation, where it also only has access to the left-hand context of each [RET] token (see section 3.3, particularly the definition of $\hat{c}_j$ following Eq. 3.2). The citing context is grouped into separate datasets based on which section




Figure 3.1.: Overview of the RAG architecture initially proposed by [LPP⁺20]. For finding the top-k documents Z' , the authors used Maximum Inner Product Search (MIPS).

in the source document it is drawn from. This results in four datasets comprised of context from the sections `IntroductionAndRelatedWork` (4.4M samples), `Methods` (124k samples), `Experiments` (497k samples) and `Conclusion` (127k samples). The candidate sections in the HDT dataset for each of these datasets are determined via exact token matching, e.g. if the lowercase title of a section contains the string "methods", it is considered a methods section and therefor used to extract samples for the `Methods` evaluation dataset. Note that the dataset `IntroductionAndRelatedWork` combines two kinds of sections commonly found in scientific documents, this is done to imitate the `ScholarCopilot` evaluation dataset that includes only these sections. `ScholarCopilot` should yield the best performance on this dataset since the model is trained exclusively on titles, abstracts, introductions and related work sections [WMN⁺25].

The dataset for the generation evaluation (section 4.3) also uses the HDT dataset. Specifically, each of the 347k samples in the dataset consists of the title, abstract and the first two and a half sentences of the introduction (the HDT dataset already provides sentences, the half-sentence is obtained by tokenizing the source sentence using the `ScholarCopilot` tokenizer and truncating it to its left-hand half based its total number of tokens). This setup mimics the examples for initial generation contexts¹ provided with the official `ScholarCopilot` implementation. Before being used for evaluation, all samples are sorted by the arXiv id of the paper they were obtained from, which effectively also sorts them by their month of publication (see [sta25b]), and only the 100k most recently published papers are retained. This is done to ensure that `ScholarCopilot` predominantly cites papers published after the paper it is prompted to generate (that is, the paper that is used as the initial generation context).

¹https://github.com/TIGER-AI-Lab/ScholarCopilot/tree/main/run_demo/src

3.3. ScholarCopilot

Instead of treating retrieval and generation as separate processing stages (see section 1) ScholarCopilot [WMN⁺25] integrates the information retrieval directly into the generation process. The model is trained to generate *retrieval tokens* ([RET]), that upon being generated, cause the model to interrupt the generation process to retrieve a reference from the retrieval corpus [WMN⁺25]. The ScholarCopilot retrieval corpus² consists of abstracts of scientific papers, which are encoded by the ScholarCopilot model prior to training/inference and stored as a HNSW [MY18] index using Faiss [DGD⁺24] for efficient retrieval. Standard ScholarCopilot uses the retrieved abstract to inform its subsequent generation by replacing the [RET] tokens in the generated context with the corresponding abstracts.

For this thesis, I evaluated the official ScholarCopilot implementation³. In their paper, the authors state that ScholarCopilot can also integrate key excerpts⁴ [WMN⁺25] but the official implementation only uses abstracts and provides no logic for extracting these excerpts. Because of that, I further on use the term "ScholarCopilot" to denote a version of this model that represents papers by their abstracts and does not use potentially more informative key excerpts.

ScholarCopilot is intended to assist with academic writing by providing "[...] text completion and citation suggestions" (see the official demo^{5,6} for an example). In the following we will thus assume that the model is used to generate a scientific paper with citations. Let

$$Y^{\text{in}} = (y_1^{\text{in}}, y_2^{\text{in}}, \dots, y_l^{\text{in}}) \quad (3.1)$$

denote the initial context given to the model. Since we seek to generate a scientific document, this context may consist of the title and abstract of the paper the model is supposed to generate. Let furthermore

$$Y_{1:i-1}^{\text{gen}} = (y_1^{\text{gen}}, y_2^{\text{gen}}, \dots, y_{j-1}^{\text{gen}}, c_j, y_{j+1}^{\text{gen}}, \dots, y_{i-1}^{\text{gen}}) \quad (3.2)$$

denote a token sequence generated by the model where c_j signifies that a [RET] token was generated by the model at generation step j . Note that the model may generate the retrieval token at multiple steps, and that the number of generated retrieval tokens is equal to the number of citations in the generated paper. The representation $\hat{c}_j$ of a [RET] token c_j computed by the model serves as an aggregated representation of the token sequence generated up to the [RET] token (e.g. $Y_{1:j-1}^{\text{gen}}$

²<https://huggingface.co/datasets/TIGER-Lab/ScholarCopilot-Data-v1>

³See <https://github.com/TIGER-AI-Lab/ScholarCopilot>. I cloned the repository on the 13th of June 2025 and evaluated that version, which corresponds to commit f8e6de9 in the source repository. As of the time of writing this (7th of September), only the README.md was extended.

⁴In the paper it says "key excerpts", likely a typo.

⁵<https://huggingface.co/spaces/TIGER-Lab/ScholarCopilot>

⁶<https://www.youtube.com/watch?v=Q1Y7S52sWDA>

in Eq. 3.2), in the same manner as BERT [DCLT19] (and also for example CF-ViT [CLL⁺23]) uses the [cls] token as a representation of its input sequence. Thus $\hat{c}_j$ is unique for every j since it is dependent of the corresponding left-hand context.

When the ScholarCopilot generator p_G generates a retrieval token (that is, $p_G(*|Y^{\text{in}}Y_{1:i-1}^{\text{gen}})$ is maximal for the [RET] token, thus $* \leftarrow c_i$), the generation process is interrupted and instead, the representation $\hat{c}_i$ is passed to the retriever $p_R(z_i|\hat{c}_i)$ used to retrieve an encoded reference z_i corresponding to an abstract

$$Y_i^{\text{ref}} = (y_{i,1}^{\text{ref}}, y_{i,2}^{\text{ref}}, \dots, y_{i,k}^{\text{ref}}) \quad (3.3)$$

from the retrieval corpus. Before passing the newly generated context to the generator the retrieval token c_i is replaced with the corresponding abstract to obtain an augmented generation output

$$Y_{1:i}^{\text{gen}} = (y_1^{\text{gen}}, y_2^{\text{gen}}, \dots, y_{i-1}^{\text{gen}}, y_{i,1}^{\text{ref}}, y_{i,2}^{\text{ref}}, \dots, y_{i,k}^{\text{ref}}) \quad (3.4)$$

that is used to inform the subsequent generation steps with the retrieved information. The generator $p_G(y_{i+1}|Y^{\text{in}}Y_{1:i}^{\text{gen}})$ predicts the next token y_{i+1} at generation step $i+1$ based on the given initial context Y^{in} and the augmented output $Y_{1:i}^{\text{gen}}$. See Algo. 1 for a detailed overview of one generation step. Before presenting the generated sequence to the user, each inserted abstract is replaced with a token referencing the cited paper (when using latex formatting, this would for example be a citation like "\cite{vaswani2017attention}"). The user is presented with this modified sequence, as well as a list of all cited references.

3.4. Passage Retrieval

As outlined in section 3.3, ScholarCopilot (SC) represents each scientific document in the retrieval corpus with its abstract. But, depending on the context for which the model is supposed to retrieve a reference, the abstract might not be the most informative passage of a candidate paper (see section 1). Thus always using a paper's abstract as reference may add irrelevant information into the generated sequence which can negatively affect the model's reasoning capabilities [SCM⁺23].

This thesis introduces a *passage retrieval* module (PR) that is supposed to enhance the citation accuracy and generation quality of ScholarCopilot (see section 1). The module is agnostic to the RAG model, which makes it possible to use with any RAG architecture. However, this thesis only investigates how it affects ScholarCopilot. First, the module takes the top- k_{SC} documents

$$Z_i = (D_1, D_2, \dots, D_{k_{\text{SC}}}) \quad (3.5)$$

returned by the SC retriever for a previously generated sequence $Y_{1:i-1}^{\text{gen}}$ and subdivides these documents into n -sized passages to obtain a set of candidate passages P_i .

Algorithm 1: Overview of one generation step of ScholarCopilot. Here Y denotes the vocabulary (the set of all tokens), Z the retrieval corpus (all encoded abstracts), G ScholarCopilot’s decoder and $\text{getAbstract}()$ a function that returns the unencoded abstract corresponding to z_i

Input: The initial context Y^{in} and the previously generated sequence $Y_{1:i-1}^{\text{gen}}$

Output: The initial context Y^{in} and the new generated sequence $Y_{1:i}^{\text{gen}}$

```

1 Algorithm: ScholarCopilotGenerationStep( $Y^{\text{in}}, Y_{1:i-1}^{\text{gen}}$ )
2  $y_i \leftarrow \arg \max_{y_i \in Y} p_G(y_i | Y^{\text{in}} Y_{1:i-1}^{\text{gen}})$ 
3 if  $y_i = [\text{RET}]$  then
4    $\hat{c}_i \leftarrow G(y_i)$ 
5    $z_i \leftarrow \arg \max_{z_i \in Z} p_R(z_i | \hat{c}_i)$ 
6    $Y_i^{\text{ref}} \leftarrow \text{getAbstract}(z_i)$ 
7    $Y_{1:i}^{\text{gen}} \leftarrow Y_{1:i-1}^{\text{gen}} Y_i^{\text{ref}}$ 
8   return  $Y^{\text{in}}, Y_{1:i}^{\text{gen}}$ 
9 end
10 else
11    $Y_{1:i}^{\text{gen}} \leftarrow Y_{1:i-1}^{\text{gen}} y_i$ 
12   return  $Y^{\text{in}}, Y_{1:i}^{\text{gen}}$ 
13 end

```

Prior work on understanding and utilizing discourse structure of prose has shown that the grammatical structure across multiple sentences can inform the reader of what is important in a document [Rum75] and can guide comprehension and recall [Man78]. Since the PR retrieval corpus (see section 3.2) consists of documents published on arXiv⁷ and thus they are curated and "[checked] for scholarly value" [sta25a], it is reasonable to assume that these documents are already well-formed and have meaningful semantic structure spanning multiple sentences and paragraphs. To preserve this structure of each document D_i in the passages obtained from it as much as possible, each document D_i is first split into sections

$$D_i = (\mathcal{E}_1, \mathcal{E}_2, \dots, \mathcal{E}_{|D_i|}) \quad (3.6)$$

by utilizing the fact that the HDT dataset [HFB⁺24] (see also section 3.2) already provides its documents structured into sections. Each section $\mathcal{E}_i$ is split into passages

$$\mathcal{E}_i = (\mathcal{P}_1, \mathcal{P}_2, \dots, \mathcal{P}_{|\mathcal{E}_i|}) \quad (3.7)$$

which are sequences of n or less tokens $\mathcal{P}_i = (t_1, t_2, \dots, t_n)$ ($\mathcal{P}_{|\mathcal{E}_i|}$ might be shorter than n tokens since $|\mathcal{E}_i| \equiv 0 \pmod n$ likely does not hold for every section in every

⁷<https://arxiv.org/>

document). This ensures that each passage only consists of sentences taken from a single section in the source document. All passages are collected into a corpus P_i and then ranked by a decoder model R_{PR} based on how relevant they are to the context $Y_{1:i-1}^{\text{gen}}$ (assuming y_i denotes the [RET] token that triggered the retrieval, see section 3.3, particularly Algo. 1). The passage retrieval prompt that is used to instruct the decoder is provided in section A.2. To achieve reliably reproducible relevance assessment performance the scores computed by R_{PR} are aggregated using Batched Self-Consistency [KDSR25]. Furthermore, to reduce the computational burden, the context is truncated to size w – yielding $Y_{i-w-1:i-1}^{\text{gen}}$ as the new context – before being passed to R_{PR} . I found that passing longform context tends to worsen the performance of the passage ranking model. The top-ranked passage is then integrated into the generated text in the same manner as the abstract would have been integrated when using standard ScholarCopilot (see Eq. 3.4). How the PR module is integrated into ScholarCopilot is shown in Algo. 2.

It is important to point out that the performance of the PR module is partially dependent on the performance of ScholarCopilot, since only the top- k_{SC} documents of ScholarCopilot are passed to the PR module (for the experiments I use $k_{SC} = 2k$ when evaluating with Recall@ k). For Recall@10 for example, it is impossible for the correct document to be chosen by the PR module if it is not in the top 20 documents ranked by ScholarCopilot. This effect can be mitigated by choosing a larger k_{SC} , but this would also drastically increase the effective runtime of the PR module, since the module computes a score for each of the p passages obtained from each of the k_{SC} documents m times (where m is the number of LLM calls for the Batched Self-Consistency score aggregation, see [KDSR25]).

The decision to let the ScholarCopilot retriever first retrieve abstracts and not directly passages is motivated by the success of employing mixed-resolution tokenization [CLL⁺23, HRBB23] for image processing with the Vision Transformer [DBK⁺20]. In [CLL⁺23, HRBB23], mixed-scale tokenization is primarily used to decrease the length of the input sequence to alleviate the computational burden imposed by the Transformer backbone whose attention layer has $O(n^2)$ time complexity (see [VSP⁺17]). Both approaches first split a given image into larger *coarse-scale* patches and then identify the most informative k patches which are then further subdivided into four *fine-scale* tokens before the image, now represented by a mixture of coarse and fine-scale tokens, is passed to the ViT backbone. ScholarCopilot and the PR retriever use the same backbone architecture (see section 4) for retrieval, so if we consider the abstract of a document D_i its coarse-scale representation and its $|D_i|$ passages (for example, when assuming 8192 tokens per document and a passage length of 512 this would mean $|D_i| = 16$) to be the fine-scale representation, then the PR module implements a similar two-stage *mixed-scale retrieval*. Let C denote the retrieval corpus and for simplicity let's assume that all documents $D_i \in C$ have the same number of $|D_i|$ passages. Then performing passage retrieval on the set of all passages would mean comparing one query to $|C| \cdot |D_i|$ passages. Using mixed-scale retrieval would result in $|C| + k_{SC}|D_i|$ comparisons, a lot less when k_{SC} is chosen

significantly smaller than $|C|$ (which is reasonable to assume and is also the case for my experiments, see section A.1). This effect is further amplified when using Batched Self-Consistency.

The abstract of a scientific document typically presents the reader with knowledge of its content, can arouse interest and can indicate whether the fulltext or part of it merits further attention [SM90]. It can serve as a condensed representation of the whole document [CO06]. An abstract can thus also be considered a semantic coarse-scale representation of a document since it usually contains short general statements about the topics discussed in the document. For example, the sentence "We propose a new simple network architecture, the Transformer, based solely on attention mechanisms, dispensing with recurrence and convolutions entirely." from the abstract of [VSP⁺17] conveys the information that the Transformer is based on attention but does not elaborate on how attention works.

On the other hand, the passages obtained from a document's sections often contain more specific information. Consider for example "We call our particular attention 'Scaled Dot-Product Attention' [...]. The input consists of queries and keys of dimension [...]" (the beginning of section 3.2.1 in [VSP⁺17]) which leads into a detailed explanation of the attention layer. Passages can thus similarly be considered *fine-scale* in a semantic sense, which allows the PR retriever to better distinguish between different aspects of the topic discussed in each paper than ScholarCopilot (that relies on abstracts alone). This could potentially result in increased retrieval performance, especially if the PR retriever is specifically trained with a contrastive loss function to identify relevant passages. However, this remains an assumption for now because training the retriever is beyond the scope of this thesis and is thus left for future research (see also section 5).

Algorithm 2: Overview of one generation step of ScholarCopilot with Passage Retrieval. `getDocument()` denotes a function that returns the fulltext paper corresponding to z_j

Input: The initial context Y^{in} and the previously generated sequence $Y_{1:i-1}^{\text{gen}}$

Output: The initial context Y^{in} and the new generated sequence $Y_{1:i}^{\text{gen}}$

```

1 Algorithm: ScholarCopilotWithPRGenerationStep( $Y^{\text{in}}, Y_{1:i-1}^{\text{gen}}$ )
2  $y_i \leftarrow \arg \max_{y_i \in Y} p_G(y_i | Y^{\text{in}} Y_{1:i-1}^{\text{gen}})$ 
3 if  $y_i = [\text{RET}]$  then
4    $\hat{c}_i \leftarrow G(y_i)$ 
5    $Z_i \leftarrow \emptyset$ 
6    $Z' \leftarrow Z$ 
7   while  $|Z_i| < k_{\text{SC}}$  do
8      $z_j \leftarrow \arg \max_{z_j \in Z'} p_R(z_j | \hat{c}_i)$ 
9      $Z_i \leftarrow Z_i \cup \{z_j\}$ 
10     $Z' \leftarrow Z' \setminus \{z_j\}$ 
11  end
12   $P_i \leftarrow \emptyset$ 
13  foreach  $z_j \in Z_i$  do
14     $D_j \leftarrow \text{getDocument}(z_j)$ 
15    foreach  $\mathcal{P}_k \in D_j$  do
16       $P_i \leftarrow P_i \cup \mathcal{P}_k$  // see Eq. 3.6 and Eq. 3.7
17    end
18  end
19   $Y_i^{\text{ref}} \leftarrow \arg \max_{\mathcal{P}_k \in P_i} p_{\text{PR}}(\mathcal{P}_k | Y_{i-w-1:i-1}^{\text{gen}})$ 
20   $Y_{1:i}^{\text{gen}} \leftarrow Y_{1:i-1}^{\text{gen}} Y_i^{\text{ref}}$ 
21  return  $Y^{\text{in}}, Y_{1:i}^{\text{gen}}$ 
22 end
23 else
24    $Y_{1:i}^{\text{gen}} \leftarrow Y_{1:i-1}^{\text{gen}} y_i$ 
25   return  $Y^{\text{in}}, Y_{1:i}^{\text{gen}}$ 
26 end

```

4. Experiments

TODO

4.1. Evaluation Methodology

The main focus of this thesis is to determine how the use of the passage retrieval module affects the retrieval accuracy and generation quality of ScholarCopilot. ScholarCopilot retains its original backbone model Qwen2.5-7B [YYZ⁺24]. Since the retriever of the passage retrieval module is implemented with Batched Self-Consistency, it is required to follow multiple permutations of the same instruction during inference reliably and consistently (see [KDSR25]). Instruction tuned models, that is, LLMs that are fine-tuned on (instruction, desired-output) pairs, are more successful at following instructions closely rather than just minimizing the next token prediction error [ZDL⁺25]. To make use of this, the PR module uses the instruction-tuned variant of the Qwen2.5 model, Qwen2.5-7B-Instruct [YYZ⁺24], as the passage retriever. Even though it has been shown that instruction tuned models may not fully utilize given instructions and are instead just good at learning the specified output format [KP23], this ability alone is valuable enough for the automated parsing of an LLMs' response to justify using an instruction tuned model (especially when using Batched Self-Consistency which greatly increases the number of responses that need to be parsed).

With ScholarCopilot as baseline, I evaluated the impact of using the PR module on the retrieval performance for a variety of retrieval contexts. To this end, I introduce the *single-step retrieval* task which aims to assess the models accuracy when resolving a single [RET] token during generation. This task closely follows the evaluation methodology used by [WMN⁺25]. For each of the section groups (IntroductionAndRelatedWork, Methods, Experiments, and Conclusion), 1k randomly selected samples are used for the evaluation. Single-step retrieval is evaluated using the Recall@*k* metric, which is given as the fraction of cases where the ground truth document appears in the top-*k* retrieved documents [WMN⁺25]. As outlined in section 3.2, the model is given the left-hand context of a citation drawn from an existing paper, together with the cited document as the ground truth. A retrieval token is appended to this context before the resulting sequence is passed to ScholarCopilot, causing the model to compute a representation of the [RET] token based on the given context. This representation is then used to compute a distribution over

the ScholarCopilot retrieval corpus that measures the relevance of each document in the corpus to the given context. When evaluating standard ScholarCopilot, the k most relevant documents according to this distribution are compared to the ground truth, otherwise the top- k_{SC} most relevant documents are passed to the passage retrieval module (for the reported results I use $k_{SC} = 2k$, see also section A.1). For each of the top- k ranked passages the PR module returns the document from which each passage was taken. These k documents are then compared to the provided ground truth. Additionally, I evaluated two different strategies for replacing other citations that appear in the citing context. Other citations are either (1) masked using a [MASK] token, this is also how [WMN⁺25] handle citations, or (2) replaced with the abstract of the cited paper, which is done to imitate the generated context the retriever would be given during inference when the model encounters a [RET] token. The datasets that use these replacement strategies are denoted as masked and with-abstracts respectively. Following the setup in [WMN⁺25], the given citing context is not truncated when passed to ScholarCopilot.¹ However, I found that passing the entire context to the PR module reduces performance, thus only the last w tokens of the citing context are used when the PR retriever is prompted to rank a given set of passages. See section A.2 for the detailed passage retrieval instruction.

Since using human experts to evaluate the quality of LLM answers is costly and slow [HPS⁺25], the quality of the generated academic text is evaluated using DeepSeek-R1-Distill-Qwen-7B [DAGY⁺25]. Following the setup in [WMN⁺25], the evaluating model (judge) is instructed to compare academic text generated by either ScholarCopilot or ScholarCopilot with passage retrieval to the ground truth document (which is the beginning of a scientific paper from arXiv, see section 3.2). It is instructed to evaluate the given generated text according to the five dimensions Content Relevance, Logical Coherence, Academic Rigor, Background Completeness and Innovation Statement, see section A.3 for the detailed instruction.² To assess the reliability of the judge, the evaluation is performed two times, once by (1) instructing the judge to only provide a score from 0 to 5 for each evaluation dimension and (2) by additionally prompting the judge to provide reasoning for its decision before assigning each score. To achieve consistent and reliable scoring, the scores for the evaluation dimensions are aggregated using Batched Self-Consistency. Specifically, the order in which the individual evaluation dimensions appear in the generation evaluation prompt is randomized every time the judge is instructed to compute a ranking and for each generated paper the judge is queried to evaluate it m_{judge} times. The evaluation is performed on 100 ground truth samples which are selected randomly from the generation evaluation dataset. To save time and computational

¹[WMN⁺25] only use the last sentence of the context when evaluating their baseline models (BM25 and E5-Mistral-7B-Instruct) but use the full context when evaluating ScholarCopilot.

²It is intentional that I don't go into detail on how these categories should be interpreted because this information is also not given in the judge instruction. How the judge interprets the evaluation categories is thus entirely dependent on the pre-training and fine-tuning data of Qwen2.5 and Deepseek-R1 respectively. This design follows [WMN⁺25] who also do not provide their judge with interpretation suggestions.

Model	Dataset	intro+relwork	methods	experiments	conclusion	Σ
SC	masked	-0.06	-0.04	-0.09	-0.07	-0.26
SC	with-abstracts	-0.05	-0.06	-0.07	-0.06	-0.24
SC+PR	masked	-0.03	-0.03	-0.05	-0.05	-0.16
SC+PR	with-abstracts	-0.04	-0.04	-0.03	-0.04	-0.15

Table 4.1.: Cumulative Recall@ k difference from $k = 10$ over $k = 5$ to $k = 1$.

resources, the text generation (by the evaluated models) is terminated early if the number of generated tokens exceeds 30k (this limit includes the abstracts/passages inserted during generation).

4.2. Single-Step Retrieval

Citation Recall. The recall results of the single-step retrieval evaluation are visualized in Figure 4.1 and Figure 4.2. Using the passage retrieval module consistently worsens the performance of ScholarCopilot (on average -0.035% / -0.027% recall on the masked/with-abstracts datasets). However, as shown in Table 4.1, when using the PR module the accumulated loss in recall from $k = 10$ to $k = 1$ is much lower compared to only using ScholarCopilot. So if given the correct document, the PR module tends to push the correct document further towards the top of the relevance ranking in comparison to the ranking computed by ScholarCopilot. This suggests that utilizing passage retrieval can improve scientific citation retrieval compared to doing retrieval solely based on abstracts.

Reproducibility. ScholarCopilot alone is unable to reproduce the retrieval performance reported in [WMN⁺25] (which is 64.8%/40.1% for Recall@10 and Recall@1 respectively). The authors do not explicitly state this in their paper but it seems they only used Introduction and Related Work sections to obtain the citing contexts for their evaluation dataset.³ Thus, using the IntroductionAndRelatedWork dataset with masked citations (see Figure 4.1 on the left) as the evaluation setup replicates their methodology (see [WMN⁺25, 4.2]) and should yield a performance similar to what is stated in [WMN⁺25]. A possible cause for the lower recall on my dataset may be that ScholarCopilot overfits to their dataset. What exactly is responsible for this significant deviation in performance is left for future research.

Recall Performance Analysis. ScholarCopilot does not yield the best performance on contexts from Introduction and Related Work sections (which is what it was trained on) but instead shows significantly better performance on Experiment sections. This

³The file scholar_copilot_eval_data_1k.json which is downloaded from <https://huggingface.co/datasets/ubowang/ScholarCopilot-TrainingData> when executing the official ScholarCopilot implementation – and I assume is what they used for their evaluation – only includes Introduction and Related Work sections.




Figure 4.1.: Comparison of single-step retrieval performance between ScholarCopilot (SC) with and without passage retrieval (PR) on the masked dataset. `intro+relwork` denotes context drawn from Introduction and Related Work sections.

is possibly due to the fact that Experiments sections often include citations in a form similar to for example "Documents are tokenized with the same byte-pair encoding as GPT-2 (Radford et al., 2019)."⁴ where the citation is directly preceded by the name of a model, specific algorithm or otherwise meaningful entity and the cited document is the paper where this algorithm was first proposed. This makes it easier for the model to identify relevant documents: taking again the previous example, a document that is about GPT-2 and would thus be a good candidate to cite after a sequence like "Documents are tokenized with the same byte-pair encoding as GPT-2" will likely frequently use the term "GPT-2". ScholarCopilot in particular benefits from this since it only retrieves from a corpus of abstracts and an abstract typically includes the name of the algorithm/method proposed in a paper while

⁴This example is taken from [LLG⁺19]




Figure 4.2.: Comparison of single-step retrieval performance between ScholarCopilot (SC) with and without passage retrieval (PR) on the with-abstracts dataset.

simultaneously being less likely than for example the Introduction to name methods that are only similar to or preceding the proposed method. As shown in Table 4.2, the Experiments dataset has indeed the highest percentage of cases where the token before the citation appears in the abstract of the cited paper. Note that this metric is likely influenced by some false positives where the last token is not a semantically meaningful entity but rather just a token that appears often in scientific abstracts.

Correct Retrieval Overlap. As shown in Table 4.3, the PR module sometimes retrieves the correct document despite ScholarCopilot being previously incorrect, especially for low k (when using Recall@ k as measure for correctness). This suggests that the PR module has the capacity to further improve citation retrieval compared to only utilizing ScholarCopilot. The fact that the performance improvement is more prevalent for small k (and consequently small k_{SC} which directly influences the number of candidate passages) also indicates that the PR module may yield

	intro+relwork	methods	experiments	conclusion
last token in abstract	0.08	0.12	0.22	0.11

Table 4.2.: Fraction of cases where the last token before the citation appears in the abstract of the cited document.

better performance if given fewer passages. To illustrate the how retrieving passages instead of abstracts can be beneficial for generative academic writing, consider the following example where ScholarCopilot fails to retrieve the correct document while the PR module is able to retrieve a passage from the correct document. Note that the following examples are slightly shortened, refer to section A.4 for the full text. The given context for retrieval was taken from [QBK⁺21]:⁵

\textbf{Retrieval context}

1 Introduction Task-oriented dialog (TOD) systems have been widely deployed for popular virtual assistants, like Siri and Google Assistant. They help people with tasks such as booking restaurants and looking for a hotel by searching databases with constraints provided by users. After retrieving a result from the database, a system may continue by conducting actions like making a reservation or providing more information about receiving the result. However, there can be multiple results from the database that match the same constraints. For example, as shown in Fig.1, the system finds two available hotels at different locations when the user is asking the system to help book a hotel. This kind of ambiguity stops system from proceeding until the system finds out which result the user looks for. [...] Our work mainly focuses on improving models ability of understanding the answers by augmenting two existing task-oriented dialog datasets: MultiWOZ <|CIT-MASK|> and Schema-Guided Dataset (SGD) #ref25#. [...] To strengthen the model with the ability to handle the ambiguity, we propose to augment these two datasets with disambiguation turns, where the system provides all possible matched results and lets the user make their own decision based on the complete information. Specifically, we first extract templates from the SIMMC 2.0 dataset <|CIT-MASK|>

The citation for which the model is supposed to retrieve a reference is the last <|CIT-MASK|> token. The context immediately preceeding the citations implies that the target reference is the publication in which the SIMMC 2.0 dataset is proposed (which would be [KMGD21]). It is further implied that this dataset is commonly used in the context of either search query disambiguation or task-oriented dialog. Given this context, ScholarCopilot retrieved the following abstract from [LH21]:

⁵"#refXX#" are citation markers from the HDT dataset the preprocessing failed to mask properly. However, the markers do not carry any meaningful information like authors, publication date or part of the paper title which they typically would include in standard LaTeX source code. Thus they are effectively already "masked" and this bug should have no significant impact on the model performance.

\textbf{Retrieved reference (SC)}

Abstract This paper presents our work on the Situated Interactive MultiModal Conversations 2.0 challenge held at Dialog State Tracking Challenge 10. SIMMC 2.0 includes 4 subtasks, and we introduce our multimodal approaches for the subtask #1, #2 and the generation of subtask #4. SIMMC 2.0 dataset is a multimodal dataset containing image and text information, which is more challenging than the problem of only text-based conversations because it must be solved by understanding the relationship between image and text. Therefore, since there is a limit to solving only text models such as BERT or GPT2, we propose a multimodal model combining image and text. We first pretrain the multimodal model to understand the relationship between image and text, then finetune our model for each task. [...]

While this abstract, and consequently the paper it is drawn from, is indeed closely linked to the SIMMC 2.0 dataset, it is not the paper that introduced this dataset but instead proposes a method for solving one of the tasks formulated in the SIMMC 2.0 paper (see [LH21, KMGD21]). While the retrieved abstract does even briefly explain what this dataset is, it also contains a lot of information that is irrelevant for the given citing context (such as what models the authors use). ScholarCopilot fails to realize that [LH21] do not propose the SIMMC 2.0 dataset themselves but rather only make use of it. A good reference for the given citing context should not only explain what the SIMMC 2.0 dataset is but also contain as little unrelated information as possible. The passage retrieval module on the other hand manages to retrieve a passage from the correct document [KMGD21]:

\textbf{Retrieved reference (SC+PR)}

The aim of the SIMMC 2.0 dataset is to emulate futuristic, real-world shopping scenarios where humans converse with a dialog agent in natural language grounded in a situated multimodal context. As a step towards this intelligent conversational agent, we leverage the dialogs and annotations in our dataset and propose four benchmark tasks (summarized in Tab.4) along with evaluation metrics. These tasks capture several multimodal conversational reasoning challenges, as elaborated next. In a real-world conversation, humans often use shorthands (coreferences) in order to refer to objects / events that have already been mentioned in the dialog. While we reserve modeling coreference resolution as a challenging task in Sec.4.2, it is important for the system to recognize ambiguous uses of such coreferences even before attempting to resolve them. [...] The multimodal disambiguation task tests this ability of the agent. More concretely, given the dialog history and the current user utterance, multimodal disambiguation requires the agent to predict a binary label conditioned on the multimodal context, to indicate the presence of a referential ambiguity in the user utterance. [...] These mentions can be resolved through (1) the dialog context

Dataset	Metric	intro+relwork	methods	experiments	conclusion	μ
masked	Recall@10	0.91	0.97	0.98	0.91	0.943
masked	Recall@5	0.94	0.94	0.96	0.94	0.945
masked	Recall@1	0.79	0.89	0.88	0.88	0.86
with-abstracts	Recall@10	0.9	0.93	0.96	0.95	0.935
with-abstracts	Recall@5	0.86	0.97	0.95	0.89	0.918
with-abstracts	Recall@1	0.87	0.87	0.93	0.84	0.878

Table 4.3.: Fraction of correct retrievals by ScholarCopilot given correct retrieval by the PR module (such that the value would be 1 if SC was correct every time the PR module was correct and 0 if SC was incorrect whenever the PR module was correct). A retrieval result is considered correct, if it satisfies the given metric.

(e.g. A: This shirt comes in XL and is \$29. #eq50# U: Please add it to cart., or (2) the multimodal context (e.g. U: How much is that red shirt?), or (3) both (e.g. U: How much is the one next to the one you mentioned?). [...])

The retrieved reference is a passage from the section that introduces the benchmark tasks proposed together with the dataset (see [KMGD21, 4]). Similar to the abstract ScholarCopilot selected as the reference for the given citing context, this passage includes information that is irrelevant (like cross-references to other parts of the paper) or too elaborate for a citation (e.g. the examples at the end), it does however provide additional relevant information. In particular, it states that SIMMC 2.0 "[emulates] futuristic, real-world shopping scenarios where humans converse with a dialog agent in natural language" which provides a possible explanation for why this dataset is useful for developing "task-oriented dialog [...] systems" but also mentions that one of the benchmark tasks is about query disambiguation, a problem that is also discussed in detail in the citing context. While the task of obtaining a concise reference by extracting only the most relevant information from a candidate passage is left for future research, this example does still highlight the potential that fine-grained passage retrieval has to identify more relevant references than what can be retrieved solely from a corpus of abstracts.

4.3. Generation Quality

Quality Comparison. The results designed to primarily assess the generation quality are shown in Figure 4.3. ScholarCopilot achieves higher scores in all evaluation categories, with a score of 15.5/25 across all dimensions. Using passage retrieval only yields a total score of 14.9/25 (−0.6), falling behind most in Background Completeness and Academic Rigor. This suggests that passage retrieval is not




Figure 4.3.: Comparison of the generation quality scores between ScholarCopilot (SC) with and without passage retrieval (PR) when the judge is instructed to provide reasoning for its scoring.

suited to improve the generation quality of academic text generation with citations.

Evaluation Reliability. Instructing the judge to provide reasoning for its generation quality assessment yields ranking scores with a mean difference of 0.11 between the scores achieved with versus without passage retrieval. Using a prompt that does not instruct the model to reason about its decision results in a much smaller mean difference of 0.01 (−0.1, see also Figure 4.4). This makes it more difficult to identify key strengths and weaknesses of the evaluated methods. Furthermore, as shown in Figure 4.5 the standard deviation of the individual ranking scores for each evaluation dimension when instructing the judge m_{judge} times to rank the same generated sequence is on average significantly higher (2.35 compared to 2.14 (−0.21) with reasoning). These results indicate that the judge is more uncertain in its evaluation when not forced to reason about its scoring. It is also important to note that the scores computed by the judge don’t always align well with human judgement, an example of this is shown in section A.5.

4.4. Additional Experiments

TODO




Figure 4.4.: Comparison of the generation quality scores between ScholarCopilot (SC) with and without passage retrieval (PR) when the judge is instructed to only provide scores.




Figure 4.5.: Comparison of the standard deviation of the computed ranking scores for each evaluation dimension when the judge is instructed with versus without asking to provide reasoning.

5. Limitations and Future Work

TODO

6. Conclusion

TODO

A. Appendix

A.1. Hyperparameters

TODO

A.2. Passage Retrieval Instruction

TODO

A.3. Generation Evaluation Instruction

TODO

A.4. Complete Example of Single-Step Retrieval

Shown below are the full context and retrieved references for the example that was used in section 4.2:

`\textbf{Retrieval context}`

1 Introduction Task-oriented dialog (TOD) systems have been widely deployed for popular virtual assistants, like Siri and Google Assistant. They help people with tasks such as booking restaurants and looking for a hotel by searching databases with constraints provided by users. After retrieving a result from the database, a system may continue by conducting actions like making a reservation or providing more information about receiving the result. However, there can be multiple results from the database that match the same constraints. For example, as shown in Fig.1, the system finds two available hotels at different locations when the user is asking the system to help book a hotel. This kind of ambiguity stops system from proceeding until the system finds out which result the user looks for. Therefore, we need to enhance the system with the ability to resolve such ambiguity brought out by multiple items returned from database search. We call this type of ambiguity as database search result ambiguity (DSR-ambiguity). Different from semantic ambiguous words (e.g. orange can be referred as

either color or fruit), the DSR-ambiguity focuses on results from multiple database search results. Solving such disambiguation tasks consists of two steps: asking clarification questions and understanding users corresponding answers. While there is a relatively larger body of literature focusing on when and how to give out the clarification question #ref23# #ref22# #ref13#, the focus on understanding users answers/intents has been relatively sparse. Our work mainly focuses on improving models ability of understanding the answers by augmenting two existing task-oriented dialog datasets: MultiWOZ <|CIT-MASK|> and Schema-Guided Dataset (SGD) #ref25#. MultiWOZ and SGD are the most popular large-scale task-oriented dialog datasets, based on which most of the state-of-the-art dialog system models are commonly trained and evaluated. According to our analysis, there are around 66[percent] dialogs of the dataset contains multiple dataset-searching results, which means the DSR-ambiguity exists. In this setting, ambiguities are skipped and the model trained based on these datasets can hardly handle the cases where users prefer to make their own choices among all the results satisfies the constraints. Furthermore, users should be given more detailed information about search results. Ideally, dialog models should provide the information and assist users to make choices, rather than picking one from the result list and recommending it to users. It is not necessary to list all the results, but enumerating 2 or 3 options would help increase users engagement. To strengthen the model with the ability to handle the ambiguity, we propose to augment these two datasets with disambiguation turns, where the system provides all possible matched results and lets the user make their own decision based on the complete information. Specifically, we first extract templates from the SIMMC 2.0 dataset <|CIT-MASK|>

\textbf{Retrieved reference (SC)}

Abstract This paper presents our work on the Situated Interactive MultiModal Conversations 2.0 challenge held at Dialog State Tracking Challenge 10. SIMMC 2.0 includes 4 subtasks, and we introduce our multimodal approaches for the subtask #1, #2 and the generation of subtask #4. SIMMC 2.0 dataset is a multimodal dataset containing image and text information, which is more challenging than the problem of only text-based conversations because it must be solved by understanding the relationship between image and text. Therefore, since there is a limit to solving only text models such as BERT or GPT2, we propose a multimodal model combining image and text. We first pretrain the multimodal model to understand the relationship between image and text, then finetune our model for each task. We achieve the 3rd best performance in subtask #1, #2 and a runner-up in the generation of subtask #4. The source code is available at <https://github.com/rungjoo/simmc2.0>.

\textbf{Retrieved reference (SC+PR)}

A.5. "False Positive" Example of LLM-based Generation Quality Evaluation

The aim of the SIMMC 2.0 dataset is to emulate futuristic, real-world shopping scenarios where humans converse with a dialog agent in natural language grounded in a situated multimodal context. As a step towards this intelligent conversational agent, we leverage the dialogs and annotations in our dataset and propose four benchmark tasks (summarized in Tab.4) along with evaluation metrics. These tasks capture several multimodal conversational reasoning challenges, as elaborated next. In a real-world conversation, humans often use shorthands (coreferences) in order to refer to objects / events that have already been mentioned in the dialog. While we reserve modeling coreference resolution as a challenging task in Sec.4.2, it is important for the system to recognize ambiguous uses of such coreferences even before attempting to resolve them. For example, consider A: The blue trousers are priced at \$45. U: What about those?, where the phrase those could be ambiguous in the following situations: (a) The user refers to a group of trousers without specifying the exact one they have in mind, (b) The user incorrectly uses a shorthand for a novel pair of trousers not mentioned in the dialog due to conversational brevity. In either cases, identifying the need for disambiguation and responding with Which ones are you talking about? The red or the green pair? is a desirable trait for a robust assistant system. The multimodal disambiguation task tests this ability of the agent. More concretely, given the dialog history and the current user utterance, multimodal disambiguation requires the agent to predict a binary label conditioned on the multimodal context, to indicate the presence of a referential ambiguity in the user utterance. This label could also be useful for other downstream tasks like assistant response generation (Sec.4.4) in order to continue the conversation in a meaningful way. We use accuracy to measure and compare model performances for this task. For this task, we aim to resolve referential mentions in user utterances to their canonical object IDs as defined for each scene. These mentions can be resolved through (1) the dialog context (e.g. A: This shirt comes in XL and is \$29. #eq50# U: Please add it to cart., or (2) the multimodal context (e.g. U: How much is that red shirt?), or (3) both (e.g. U: How much is the one next to the one you mentioned?). The input for

Note that "%" is replaced with [percent] due to formatting issues.

A.5. "False Positive" Example of LLM-based Generation Quality Evaluation

TODO

Part of the Introduction generated by SC

[...]

In this paper, we propose a reversible adversarial example generator (RAEG) that introduces slight changes to the images to fool traditional

classification models. Even though malicious attacks train pirated models based on the defended versions of the protected images, RAEG can significantly weaken the functionality of these models. Meanwhile, the reversibility of RAEG ensures the performance of authorized models. Therefore, the protected data can be used for training the authorized models, but not for pirated models.\n\nThe second problem is how to defend against adversarial defense methods. Adversarial defense methods aim to eliminate the adversarial perturbations from adversarial examples to recover the original data. These methods can be divided into two categories: adversarial training \\\cite{arxivID-1705.03387-0, arxivID-2310.05947-1, arxivID-1905.11713-2, arxivID-2109.09901-3, arxivID-1906.11567-4, arxivID-1711.08001-5, arxivID-2007.14672-6, arxivID-2003.02484-7} and adversarial purification \\\cite{arxivID-1912.05391-8, arxivID-1612.07767-9, arxivID-1804.01635-10, arxivID-2405.15971-11, arxivID-2405.15971-12, arxivID-2212.05630-13, arxivID-2406.03143-14}.

[...]

Part of the Introduction generated by SC+PR

[...]

We develop a reversible adversarial example generator (RAEG) that introduces slight changes to the images to fool traditional classification models. The changes are imperceptible to the human eyes. The protected images cannot be used to train models directly. However, the protected images can be reversely transformed into the original images. The original images can be used to train models and the authorized models can achieve good performance. Meanwhile, the protected images cannot be used to train pirated models. The pirated models trained based on the protected images cannot achieve satisfactory performance. The pipeline of our proposed method is shown in Figure~\\ref{fig1}. The protected images can be used to train models and the authorized models can achieve good performance. Meanwhile, the protected images cannot be used to train pirated models. The pirated models trained based on the protected images cannot achieve satisfactory performance.

[...]

Adversarial Training: Adversarial training is a proactive defense method. Adversarial training is used to train models that are robust to adversarial attacks. Adversarial training is a widely used defense method. Adversarial training is used to train models that are robust to adversarial attacks. Adversarial training is a widely used defense method. Adversarial training is used to train models that are robust to adversarial attacks. Adversarial training is a widely used defense method.

[...]

Scores computed by the judge

A.5. "False Positive" Example of LLM-based Generation Quality Evaluation

```
"sc_scores": {  
  "Content Relevance": 3.4,  
  "Logical Coherence": 2.8,  
  "Academic Rigor": 3.0,  
  "Background Completeness": 2.8,  
  "Innovation Statement": 2.4,  
  "total": 13.4  
},  
"sc_pr_scores": {  
  "Content Relevance": 4.2,  
  "Logical Coherence": 4.0,  
  "Academic Rigor": 3.4,  
  "Background Completeness": 3.6,  
  "Innovation Statement": 4.4,  
  "total": 18.6  
}
```

Note that getting stuck in generating repetitive sequences is not an issue introduced by the PR module, ScholarCopilot also occasionally shows this behaviour, see for example `records_eval_generation_2025-08-30_11-24-16.json`, line 2738-2759.

Bibliography

- [AHS⁺24] Akari Asai, Jacqueline He, Rulin Shao, Weijia Shi, Amanpreet Singh, Joseph Chee Chang, Kyle Lo, Luca Soldaini, Sergey Feldman, Mike D’arcy, David Wadden, Matt Latzke, Minyang Tian, Pan Ji, Shengyan Liu, Hao Tong, Bohao Wu, Yanyu Xiong, Luke Zettlemoyer, Graham Neubig, Dan Weld, Doug Downey, Wen tau Yih, Pang Wei Koh, and Hannaneh Hajishirzi. Openscholar: Synthesizing scientific literature with retrieval-augmented lms. *arXiv preprint arXiv:2411.14199*, 2024.
- [CLL⁺23] Mengzhao Chen, Mingbao Lin, Ke Li, Yunhang Shen, Yongjian Wu, Fei Chao, and Rongrong Ji. Cf-vit: A general coarse-to-fine method for vision transformer. In *Proceedings of the AAAI conference on artificial intelligence*, volume 37, pages 7042–7052, 2023.
- [CO06] Cate Cross and Charles Oppenheim. A genre analysis of scientific abstracts. *Journal of Documentation*, 62:428–446, 07 2006.
- [DAGY⁺25] DeepSeek-AI, Daya Guo, Dejian Yang, Haowei Zhang, Junxiao Song, Ruoyu Zhang, Runxin Xu, Qihao Zhu, Shirong Ma, Peiyi Wang, Xiao Bi, Xiaokang Zhang, Xingkai Yu, Yu Wu, Z. F. Wu, Zhibin Gou, Zhihong Shao, Zhuoshu Li, Ziyi Gao, Aixin Liu, Bing Xue, Bingxuan Wang, Bochao Wu, Bei Feng, Chengda Lu, Chenggang Zhao, Chengqi Deng, Chenyu Zhang, Chong Ruan, Damai Dai, Deli Chen, Dongjie Ji, Erhang Li, Fangyun Lin, Fucong Dai, Fuli Luo, Guangbo Hao, Guanting Chen, Guowei Li, H. Zhang, Han Bao, Hanwei Xu, Haocheng Wang, Honghui Ding, Huajian Xin, Huazuo Gao, Hui Qu, Hui Li, Jianzhong Guo, Jiashi Li, Jiawei Wang, Jingchang Chen, Jingyang Yuan, Junjie Qiu, Junlong Li, J. L. Cai, Jiaqi Ni, Jian Liang, Jin Chen, Kai Dong, Kai Hu, Kaige Gao, Kang Guan, Kexin Huang, Kuai Yu, Lean Wang, Lecong Zhang, Liang Zhao, Litong Wang, Liyue Zhang, Lei Xu, Leyi Xia, Mingchuan Zhang, Minghua Zhang, Minghui Tang, Meng Li, Miaojun Wang, Mingming Li, Ning Tian, Panpan Huang, Peng Zhang, Qiancheng Wang, Qinyu Chen, Qiushi Du, Ruiqi Ge, Ruisong Zhang, Ruizhe Pan, Runji Wang, R. J. Chen, R. L. Jin, Ruyi Chen, Shanghao Lu, Shangyan Zhou, Shanhuang Chen, Shengfeng Ye, Shiyu Wang, Shuiping Yu, Shunfeng Zhou, Shuting Pan, S. S. Li, Shuang Zhou, Shaoqing Wu, Shengfeng Ye, Tao Yun, Tian Pei, Tianyu Sun, T. Wang, Wangding Zeng, Wanbiao Zhao, Wen Liu, Wenfeng Liang, Wenjun Gao, Wenqin Yu, Wentao Zhang, W. L. Xiao,

Bibliography

- Wei An, Xiaodong Liu, Xiaohan Wang, Xiaokang Chen, Xiaotao Nie, Xin Cheng, Xin Liu, Xin Xie, Xingchao Liu, Xinyu Yang, Xinyuan Li, Xuecheng Su, Xuheng Lin, X. Q. Li, Xiangyue Jin, Xiaojin Shen, Xiaosha Chen, Xiaowen Sun, Xiaoxiang Wang, Xinnan Song, Xinyi Zhou, Xianzu Wang, Xinxia Shan, Y. K. Li, Y. Q. Wang, Y. X. Wei, Yang Zhang, Yanhong Xu, Yao Li, Yao Zhao, Yaofeng Sun, Yaohui Wang, Yi Yu, Yichao Zhang, Yifan Shi, Yiliang Xiong, Ying He, Yishi Piao, Yisong Wang, Yixuan Tan, Yiyang Ma, Yiyuan Liu, Yongqiang Guo, Yuan Ou, Yudian Wang, Yue Gong, Yuheng Zou, Yujia He, Yunfan Xiong, Yuxiang Luo, Yuxiang You, Yuxuan Liu, Yuyang Zhou, Y. X. Zhu, Yanhong Xu, Yanping Huang, Yaohui Li, Yi Zheng, Yuchen Zhu, Yunxian Ma, Ying Tang, Yukun Zha, Yuting Yan, Z. Z. Ren, Zehui Ren, Zhangli Sha, Zhe Fu, Zhean Xu, Zhenda Xie, Zhengyan Zhang, Zhewen Hao, Zhicheng Ma, Zhigang Yan, Zhiyu Wu, Zihui Gu, Zijia Zhu, Zijun Liu, Zilin Li, Ziwei Xie, Ziyang Song, Zizheng Pan, Zhen Huang, Zhipeng Xu, Zhongyu Zhang, and Zhen Zhang. Deepseek-r1: Incentivizing reasoning capability in llms via reinforcement learning. *arXiv preprint arXiv:2501.12948*, 2025.
- [DBK⁺20] Alexey Dosovitskiy, Lucas Beyer, Alexander Kolesnikov, Dirk Weissenborn, Xiaohua Zhai, Thomas Unterthiner, Mostafa Dehghani, Matthias Minderer, Georg Heigold, Sylvain Gelly, Jakob Uszkoreit, and Neil Houlsby. An image is worth 16x16 words: Transformers for image recognition at scale. *arXiv preprint arXiv:2010.11929*, 2020.
- [DCLT19] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. Bert: Pre-training of deep bidirectional transformers for language understanding. In *Proceedings of the 2019 conference of the North American chapter of the association for computational linguistics: human language technologies, volume 1 (long and short papers)*, pages 4171–4186, 2019.
- [DGD⁺24] Matthijs Douze, Alexandr Guzhva, Chengqi Deng, Jeff Johnson, Gergely Szilvasy, Pierre-Emmanuel Mazaré, Maria Lomeli, Lucas Hosseini, and Hervé Jégou. The faiss library. *arXiv preprint arXiv:2401.08281*, 2024.
- [DOW⁺24] Hyo Jin Do, Rachel Ostrand, Justin D Weisz, Casey Dugan, Prasanna Sattigeri, Dennis Wei, Keerthiram Murugesan, and Werner Geyer. Facilitating human-llm collaboration through factuality scores and source attributions. *arXiv preprint arXiv:2405.20434*, 2024.
- [HFB⁺24] Haoyu He, Markus Flicke, Jan Buchman, Iryna Gurevych, and Andreas Geiger. HDT: Hierarchical Document Transformer. Conference on Language Modeling, 2024.
- [HPS⁺25] Sebastian Heineking, Jonas Probst, Daniel Steinbach, Martin Potthast, and Harris Scells. Ranking generated answers: On the agreement of retrieval models with humans on consumer health questions. In

- European Conference on Information Retrieval*, pages 128–137. Springer, 2025.
- [HRBB23] Jakob Drachmann Havtorn, Amélie Royer, Tijmen Blankevoort, and Babak Ehteshami Bejnordi. Msvit: Dynamic mixed-scale tokenization for vision transformers. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 838–848, 2023.
- [KDSR25] Anton Korikov, Pan Du, Scott Sanner, and Navid Rekabsaz. Batched self-consistency improves llm relevance assessment and ranking. *arXiv preprint arXiv:2505.12570*, 2025.
- [KMGD21] Satwik Kottur, Seungwhan Moon, Alborz Geramifard, and Babak Damavandi. Simmc 2.0: A task-oriented dialog dataset for immersive multimodal conversations. *arXiv preprint arXiv:2104.08667*, 2021.
- [KOM⁺20] Vladimir Karpukhin, Barlas Oguz, Sewon Min, Patrick SH Lewis, Ledell Wu, Sergey Edunov, Danqi Chen, and Wen-tau Yih. Dense passage retrieval for open-domain question answering. In *EMNLP (1)*, pages 6769–6781, 2020.
- [KP23] Po-Nien Kung and Nanyun Peng. Do models really learn to follow instructions? an empirical study of instruction tuning. *arXiv preprint arXiv:2305.11383*, 2023.
- [LH21] Joosung Lee and Kijong Han. Multimodal interactions using pretrained unimodal models for simmc 2.0. *arXiv preprint arXiv:2112.05328*, 2021.
- [LLG⁺19] Mike Lewis, Yinhan Liu, Naman Goyal, Marjan Ghazvininejad, Abdelrahman Mohamed, Omer Levy, Ves Stoyanov, and Luke Zettlemoyer. Bart: Denoising sequence-to-sequence pre-training for natural language generation, translation, and comprehension. *arXiv preprint arXiv:1910.13461*, 2019.
- [LOS⁺23] Jakub Lála, Odhran O’Donoghue, Aleksandar Shtedritski, Sam Cox, Samuel G Rodriques, and Andrew D White. Paperqa: Retrieval-augmented generative agent for scientific research. *arXiv preprint arXiv:2312.07559*, 2023.
- [LPP⁺20] Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, Sebastian Riedel, and Douwe Kiela. Retrieval-augmented generation for knowledge-intensive nlp tasks. *Advances in neural information processing systems*, 33:9459–9474, 2020.
- [Man78] Jean M. Mandler. A code in the node: The use of a story schema in retrieval. *Discourse Processes*, 1(1):14–35, 1978.

Bibliography

- [MY18] Yu A Malkov and Dmitry A Yashunin. Efficient and robust approximate nearest neighbor search using hierarchical navigable small world graphs. *IEEE transactions on pattern analysis and machine intelligence*, 42(4):824–836, 2018.
- [QBK⁺21] Kun Qian, Ahmad Beirami, Satwik Kottur, Shahin Shayandeh, Paul Crook, Alborz Geramifard, Zhou Yu, and Chinnadhurai Sankar. Database search results disambiguation for task-oriented dialog systems. *arXiv preprint arXiv:2112.08351*, 2021.
- [RRS20] Adam Roberts, Colin Raffel, and Noam Shazeer. How much knowledge can you pack into the parameters of a language model? *arXiv preprint arXiv:2002.08910*, 2020.
- [Rum75] David E. Rumelhart. Notes on a schema for stories. In Daniel G. Bobrow and Allan Collins, editors, *Representation and Understanding*, pages 211–236. Morgan Kaufmann, San Diego, 1975.
- [SCM⁺23] Freda Shi, Xinyun Chen, Kanishka Misra, Nathan Scales, David Dohan, Ed H Chi, Nathanael Schärli, and Denny Zhou. Large language models can be easily distracted by irrelevant context. In *International Conference on Machine Learning*, pages 31210–31227. PMLR, 2023.
- [SHC⁺24] Aviv Slobodkin, Eran Hirsch, Arie Cattan, Tal Schuster, and Ido Dagan. Attribute first, then generate: Locally-attributable grounded text generation. *arXiv preprint arXiv:2403.17104*, 2024.
- [SM90] Françoise Salager-Meyer. Discoursal flaws in medical english abstracts: A genre analysis per research- and text-type. *Text - Interdisciplinary Journal for the Study of Discourse*, 10, 01 1990.
- [SQK⁺25] Rulin Shao, Rui Qiao, Varsha Kishore, Niklas Muennighoff, Xi Victoria Lin, Daniela Rus, Bryan Kian Hsiang Low, Sewon Min, Wen-tau Yih, Pang Wei Koh, and Luke Zettlemoyer. Reasonir: Training retrievers for reasoning tasks. *arXiv preprint arXiv:2504.20595*, 2025.
- [sta25a] arXiv staff. About arXiv, 2025. <https://info.arxiv.org/about/index.html>, Last accessed: 10.09.2025.
- [sta25b] arXiv staff. Understanding the arXiv identifier, 2025. https://info.arxiv.org/help/arxiv_identifier.html, Last accessed: 15.09.2025.
- [STB25] Shubham Sharma, Sneha Tuli, and Narendra Badam. Challenges and applications of large language models: A comparison of gpt and deepseek family of models. *arXiv preprint arXiv:2508.21377*, 2025.
- [VSP⁺17] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention

- is all you need. *Advances in neural information processing systems*, 30, 2017.
- [WMN⁺25] Yubo Wang, Xueguang Ma, Ping Nie, Huaye Zeng, Zhiheng Lyu, Yuxuan Zhang, Benjamin Schneider, Yi Lu, Xiang Yue, and Wenhui Chen. Scholarcopilot: Training large language models for academic writing with accurate citations. *arXiv preprint arXiv:2504.00824*, 2025.
- [XLS⁺24] Fangyuan Xu, Kyle Lo, Luca Soldaini, Bailey Kuehl, Eunsol Choi, and David Wadden. Kiwi: A dataset of knowledge-intensive writing instructions for answering research questions. *arXiv preprint arXiv:2403.03866*, 2024.
- [XWL⁺24] Sirui Xia, Xintao Wang, Jiaqing Liang, Yifei Zhang, Weikang Zhou, Jiaji Deng, Fei Yu, and Yanghua Xiao. Ground every sentence: Improving retrieval-augmented llms with interleaved reference-claim generation. *arXiv preprint arXiv:2407.01796*, 2024.
- [YYZ⁺24] An Yang, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chengyuan Li, Dayiheng Liu, Fei Huang, Haoran Wei, Huan Lin, Jian Yang, Jianhong Tu, Jianwei Zhang, Jianxin Yang, Jiaxi Yang, Jingren Zhou, Junyang Lin, Kai Dang, Keming Lu, Keqin Bao, Kexin Yang, Le Yu, Mei Li, Mingfeng Xue, Pei Zhang, Qin Zhu, Rui Men, Runji Lin, Tianhao Li, Tianyi Tang, Tingyu Xia, Xingzhang Ren, Xuancheng Ren, Yang Fan, Yang Su, Yichang Zhang, Yu Wan, Yuqiong Liu, Zeyu Cui, Zhenru Zhang, and Zihan Qiu. Qwen2 technical report. *arXiv preprint arXiv:2407.10671*, 2024.
- [ZBL⁺24] Jiajie Zhang, Yushi Bai, Xin Lv, Wanjuan Gu, Danqing Liu, Minhao Zou, Shulin Cao, Lei Hou, Yuxiao Dong, Ling Feng, and Juanzi Li. Longcite: Enabling llms to generate fine-grained citations in long-context qa. *arXiv preprint arXiv:2409.02897*, 2024.
- [ZDL⁺25] Shengyu Zhang, Linfeng Dong, Xiaoya Li, Sen Zhang, Xiaofei Sun, Shuhe Wang, Jiwei Li, Runyi Hu, Tianwei Zhang, Fei Wu, and Guoyin Wang. Instruction tuning for large language models: A survey. *arXiv preprint arXiv:2308.10792*, 2025.